

What is search for?

- single agent
- deterministic actions
- fully observed state
- discrete state space

- planning: sequence of actions

- the path to the goal is the important thing
- paths have different costs
- Heuristics give problem specific guidance

- identification: assignments to variables

- the goal is important, not the path
- all paths have same depth (for some formulations)
- CSP are a specialized kind

CSP

standard search problem:

- state is a "black box": arbitrary data structure
- powerful but limiting abstraction
- goal test and successor functions could be anything

CSP's: we have a structure

- a state is made of variables X_i with values from a domain D (sometimes D depends on i)
- goal test is a set of constraints that tell us legal combinations of values for subsets of variables

Constraint graphs

an arc for each constraint, doesn't tell which constraint it is but only that there is one

binary CSP: all constraint relates at most two variables

general purpose CSP algorithms use the graph structure to speed up search

Varieties of CSP

- discrete variables

- finite domains

size of means $O(d^n)$ complete assignments

e.g. boolean CSP

- infinite domains (integers, strings etc.)

e.g. job scheduling, variables start/end times for each job

linear constraints solvable, non linear undecidable

- continuous variables

e.g. start/end times for telescope observation

linear constraints solvable in polynomial time by LP method

Varieties of constraints

- unary constraints: involve a single variable (= reducing domains)

e.g. SA ≠ green

- binary constraints

e.g. SA ≠ WA

- Higher-order constraints

Preferences (soft constraints)

often representable by a cost for each variable assignment

gives constrained optimization problems (ignore until Bayes nets)

Solving CSPs

standard formulation

states defined by values assigned so far (partial assignments)

- initial state: empty assignment $\{\}$

- successor function: assigns value to variable

- goal state: the current assignment is complete and satisfies all constraints

↳ BFS is bad bcz it has to explore the whole tree.
we'll use DFS based approaches

Naive search problem: tries all assignments

Backtracking search

basic uninformed alg. for solving CSPs

we apply to DFS the following improvements

1) one variable at a time

- variable assignments are commutative, so fix ordering
i.e. WA=red then NT=green same as NT=green then WA=red
- each step assignment to a single variable

2) check constraints as you go

i.e. consider only values that do not conflict with previous assignments

Improving backtracking

• Ordering : what variable should we assign next?
in what order should its values be tried?

• Structure : exploit the problem structure

• filter : can we detect inevitable failure early?

Filtering: Keep track of domains for unassigned variables and cross off bad options
Forward checking

- don't pursue values that violate a constraint when added to the existing assignment
- doesn't provide early detection for all failures
- Only propagates info from assigned to unassigned variables, can't see conflicts between unassigned variables.

Arc Consistency

An arc $x \rightarrow y$ is **consistent** if and only if $\forall x$ in the tail $\exists y$ in the head which could be assigned without any violation of constraints.

To make an arc consistent we can remove elements from the tail

- if x loses a value neighbours of x must be checked again
- Arc consistency detects failure earlier than forward checking
- can be run as a preprocessor or after each assignment

downside: really runtime heavy (like A*)

An algorithm for this that only works for binary CSP is AC-3, runtime $O(n^2d^3)$

Limitations:

- can have one solution left
- multiple solutions left
- no solutions left

Ordering

- variable ordering**: minimum remaining values (MRV)
choose variable with less legal values in its domain
min bcz it will fail fast

- value ordering**: least constraining value
given a var. choose the lcv
ie. the one that rules out the fewest values in the remaining variables
it may take some computation to determine

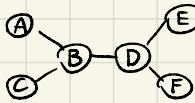
K-consistency

- z-consistency: for each pair of nodes, any consistent assignment to one can be extended to the other (the one we saw before)
- K-consistency: for each K-nodes any consistent assignment to K-1 can be extended to the K-th node
- strong K consistency: also K-1, K-2... consistent
- n-consistency means we can solve without backtracking
↑ n nodes in graph

Structure

extreme case: independent sub-problems \rightarrow identifiable as connected components of constraint graph
tree structured CSP

if the constraint graph has no loops the CSP can be solved in $O(nd^2)$ (instead of the typical $O(d^n)$)



algorithm:

• order: choose root variable, order parent \rightarrow children



• remove backward: for $i = n-2$ apply `removeInconsistent(parent(xi), xi)`
we only have one arc pointing to a node bcz it's a tree

claim: after backward pass, all root-to-leaf arcs are consistent

claim: if root-to-leaf arcs are consistent, forward assignment will not back track

Nearly tree-structured CSPs

making the structure a tree by deleting nodes

conditioning: instantiate a variable, prune its neighbor's domains
(assign)

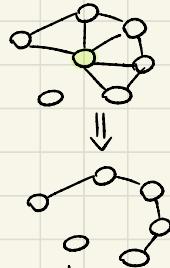
• cutset conditioning: instantiate in all possible ways a set of variables such that the remaining graph is a tree
resolve the tree and you found a solution if united with the assignments

cutset of size c runtime $O(d^c \cdot (n-c)d^2)$ very fast for very small c

tree decomposition

idea: make a tree of mega-variables

each mega variable is part of the original CSP
subproblems overlap to ensure consistent solutions



Iterative algorithms for CSPs

typically work with "complete" states, i.e. all variables assigned

N.B. not necessarily legal

to apply to CSPs:

- take an assignment with unsatisfied constraints
- operators (instead of successor) reassign variable values
- no fringe

algorithm:

· variable selection: random conflicted variable

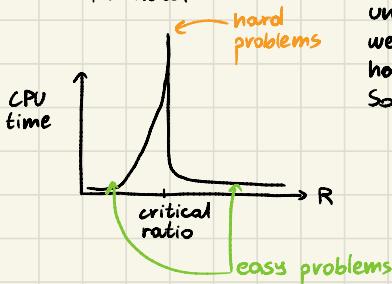
· value selection: min-conflicts heuristic

choose a value that violates the fewest constraints

performance: really fast really often

but the performance depends on the ratio:

$$R = \frac{\text{constraints}}{\text{variables}}$$



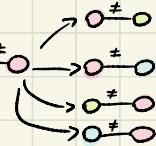
unfortunate because
we want to solve
hard problems.

So min conflicts is often ineffective in practice

Local search

- tree search: keep unexplored alternatives on the fringe (ensures completeness)
- local search: improve a single option until you can't make better (no fringe)

new successor function: local changes



generally much faster and more memory efficient (but incomplete and sub-optimal)

Hill climbing

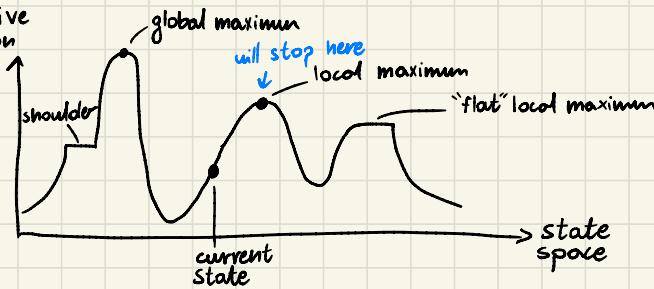
idea: start wherever

repeat: move to the best neighboring state
if no neighbors better than current, quit

not complete, not optimal

what's good? Really easy to apply

objective function



Simulated annealing

idea: Escape local maxima by allowing downhill moves

but makes them rarer as you go

it has a guaranteed optimality

