

Determine how the state of a process evolve

long term scheduling → determines the number of active processes

a new process is inactive until the long-term scheduler activates it

- when can a new process be activated?

- which process to activate?

more active processes increase the execution times of the active processes (and requires more virtual mem)

So the scheduler will monitor the idle time of the processor and activate processes when the idle time becomes too high it will activate a process.

- The type of process is also important.

- Too many I/O-bound processes can lead to all process being blocked.

- interactive processes are often allowed to be active immediately

medium term scheduling → determines which processes are (partially) present in physical memory.

non real-time short-term scheduling → which of the ready processes may use the processor

turn-around time → T_q how long to be completely executed (from when it was started)

It's desirable that the processor is almost always working. High utilisation often results in larger turn around time. So we have to make a compromise.

Fairness → capacity of the resources is fairly distributed over the processes

predictability → how sensitive the exec. time of a process is to the presence of other processes

Notation

$\alpha | \beta | \gamma$ ↗ boundary conditions
number of processes ↘
function to minimize ↙

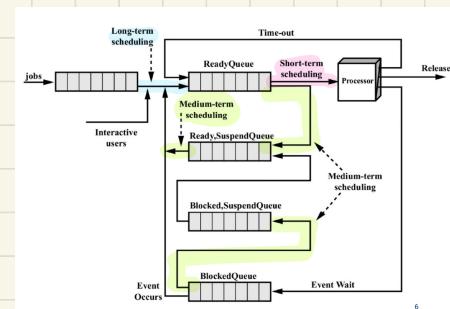
• $\alpha = 1$ 1 processor $\alpha = p_m$ m identical processors

• β : r_j → tasks with arrival times $\neq 0$
pmth → tasks can be interrupted

prec → some tasks have to wait for others to finish

• γ : C_j → completion time of task j
we either consider $\gamma = \sum C_i$ or $\gamma = \sum w_i C_i$ ← weighted sum

$\sum U_j$ number of tasks that don't meet deadline



non real time short term scheduling algorithms

tasks can be interrupted \rightarrow preemptive systems.

(FIFO)

First come first served (FCFS)

Serve the processes in the order they were placed in the queue
drawbacks:

- small tasks end up behind large processes
- I/O bound processes are disadvantaged, bcz everytime an I/O op. happens they go back to the queue

Round Robin (RR)

set a maximum time a process can have continuous access to the processor (quantum).

After this time is ended the tasks is interrupted and back to the ready queue

shorter quantum \rightarrow a lot of time spent on switching processes instead of executing them

Still disadvantage for I/O bound processes. A solution may be using a different queue for them.
the amount of time they are allowed is quantum-time they had access to the processor since the last time they were in the queue

Shortest process next (SPN)

choose the process that requires the least processor time.

optimal for $1 \mid \Sigma C_i \rightarrow$ tasks arriving together at time $t=0$, no interruptions

not for $1 \mid r_i \mid \Sigma C_i$

difficult to implement because requires prior knowledge of the running time of all processes.

Disadvantages for large processes

proof of optimality

consider Turnaround time $T_q \sim \sum C_i$ p_j processing time of proc. j

$$T_q = p_1 + (p_1 + p_2) + \dots + (p_1 + \dots + p_{i-1}) + p_i + (p_1 + \dots + p_{i+1}) + \dots + p_N = \sum_{i=1}^N (N-i+1)p_i \quad \text{optimal if } p_1 \leq p_2 \leq \dots \leq p_N$$

t has to wait for p_i to start

Weighted shortest process next (WSPN)

chooses processes according to p_j/w_j $w_j \geq 0$ is the weight

optimal $1 \mid \Sigma w_j C_j$

aims at minimizing weighted turnaround time $T_q \sim \frac{1}{N} \sum w_j C_j$

$$w_1 p_1 + w_2 (p_1 + p_2) + \dots + w_N (p_1 + \dots + p_{N-1} + p_N) = w_1 (p_1 + p_2 + \dots + p_{i-1} + p_i)$$

proof of optimality

consider order S optimal but not for WSPN, then $\frac{p_i}{w_i} > \frac{p_{i+1}}{w_{i+1}} \rightarrow p_i w_{i+1} > p_{i+1} w_i$

turnaround time for i and $i+1$

$$S = w_i (p_1 + p_2 + \dots + p_{i-1} + p_i) + w_{i+1} (p_1 + p_2 + \dots + p_{i-1} + p_i + p_{i+1})$$

now consider order S' obtained from switching processes i and $i+1$ in S

$$W_{i+1}(p_1 + p_2 + \dots + p_{i-1} + p_{i+1}) + W_i(p_1 + p_2 + \dots + p_{i-1} + p_{i+1} + p_i)$$

↓ switch for convenience

$$S = W_i(p_1 + p_2 + \dots + p_{i-1} + p_{i+1} + p_i) + W_{i+1}(p_1 + p_2 + \dots + p_{i-1} + p_{i+1})$$

Now subtract: $S - S'$

$$W_i(p_1 + p_2 + \dots + p_{i-1} + p_i) - W_i(p_1 + p_2 + \dots + p_{i-1} + p_{i+1} + p_i) = -W_i p_{i+1}$$

$$W_{i+1}(p_1 + p_2 + \dots + p_{i-1} + p_i + p_{i+1}) - W_{i+1}(p_1 + p_2 + \dots + p_{i-1} + p_{i+1}) = W_{i+1} p_i$$

$$W_{i+1} p_i - W_i p_{i+1} > 0 \rightarrow T_q \text{ decreased!} \Rightarrow S \text{ was not optimal}$$

Shortest remaining time next (SRTN)

tasks can be interrupted (assume int. itself doesn't waste time).

When new process becomes ready if it needs less running time than the remaining time of the current process, the current one gets interrupted and replaced with the new one.
optimal for optimizing turnaround time.

$$1/r_j, p_{\text{max}} / \sum C_i$$

it requires perfect knowledge of the length of all tasks, also have to track the remaining exec. time. Starvation for longer tasks

proof of optimality

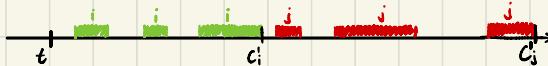
Consider order S optimal but not for SRTN, then



$$C'_j = \max(C_i, C_j)$$

$$C'_i < \min(C_i, C_j)$$

consider now another order S' so that



Compare turnaround times

$$\sum_k C_k = \sum_{k \neq i, j} C_k + C_i + C_j > \sum_{k \neq i, j} C_k + C'_i + C'_j$$

T_q has decreased

S was not optimal

Nonpreemptive SRTN (nSRTN)

tries to minimize the running time when tasks are not interrupted

$$-|r_i| \sum C_j$$

given the optimal order $C_1 < C_2 \dots C_N$

construct same sequence with times $\bar{C}_1 + \bar{C}_2 + \dots + \bar{C}_N$

possibly leave processor idle until next process arrives

its only possible to reach $2x$ of the optimal time ($2x$ approximation)

proof

prove that $\sum_j \bar{C}_j \leq 2 \sum_j C_j$

consider completion time C_j

$\bar{C}_j \leq \max_{k \leq i} r_k + \sum_{k=1}^i p_k$ maximal if the first task is the last of the j -task to arrive
(has to wait)

for $\max_{k \leq i} r_k \leq C_j \rightarrow \max_{k \leq j} r_k \leq C_j \rightarrow \sum_{k=1}^j p_k \leq C_j \rightarrow \bar{C}_j \leq 2C_j$

Since tasks are processed in order $\rightarrow \sum_{k=1}^j p_k \leq C_j$

Highest response ratio next (HRRN)

each process has a weighted turnaround time equal to $RR = (w + p_s) / p_s$ w = time process has been waiting

The increase is proportional to the processing time p_s of the project (shorter increase faster)

When a process is released we choose the one with the highest RR. Short processes advantaged.

Requires at least estimates of times

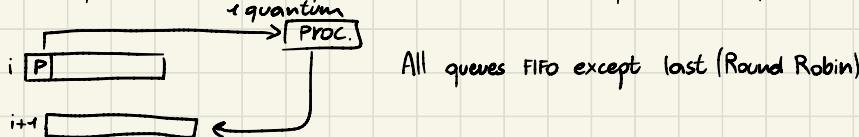
Multilevel feedback scheduling

look at how much time a process already had at its disposal

We allow processes that had access to the processor less frequently

An implementation uses n ready queues with a priority mechanism.

Process in queue i can access processor only if the $i-1$ queues are empty.



Short tasks are favored. The quantum is increased as i increases.

I/O bound processes always in queue 1

Traditional Unix Scheduler

Similar to multilevel feedback queue.

Time divided into equal intervals, each process j has a priority $P_j(i)$ during interval i .

The lower the value the higher the priority.

The scheduler resolution indicates how often a clock interrupt is generated per second (typically 60/s).

For each process j a counter $U_j(i)$ is kept, it indicates how often the process was active during i .
at each interrupt the counter of the current process increased by 1.

At the end of each interval we calculate the average CPU usage

$$CPU_j(i) = U_j(i)/z + CPU_j(i-1)/z$$

then we adjust the priorities

$$P_j(i) = Base_j + CPU_j(i-1)/z + nice_j$$

Fair-share scheduler

Add another $GU_k(i)$ per group K . Increase as before if a process from group K was active.

$$GCP\ U_k(i) = GU_k(i)/z + GCP\ U_k(i-1)/z$$

$$P_j(i) = Base_j + CPU_j(i-1)/z + GCP\ U_k(i-1)/z \cdot w_k + nice_j$$

w_k is a group weight, $\sum w_k = 1$ percentage of use of CPU for group K

Real time short time scheduling

deadline d_i for process i

Divided in static and dynamic \rightarrow online

↳ offline (taken when system not operational)

- 1) static table-driven solutions: table created in advance that perfectly represents which tasks will be executed when. If not all can meet their deadlines some processes will not be executed
- 2) static priority-driven solutions: each task is given a priority as a result of a scheduling algorithm

With dynamic solutions decisions regard the acceptance of new tasks, this can be done "best effort"

Not all tasks actually have a deadline.

Some deadlines may be hard (never to be crossed) or soft (small crossing is solvable)

Deadline Scheduling

Assume we know all processes, their duration p_i and deadline d_i .

$C_j \rightarrow$ completion time of j

$U_j = 1$ if $C_j > d_j$ task j not completed in time, otherwise $U_j = 0$

We are interested in the number of tasks that fail the deadline ($\sum U_i$)

or the maximum degree to which a task is late ($L_{\max} = \max_{i=1}^n (C_i - d_i)$)
those add up to γ (function we wish to minimize)

Earliest deadline first (EDF) - Earliest Due Date (EDD)

Schedule tasks according to ascending deadline

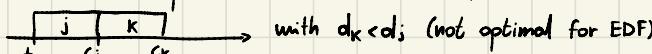
task i has deadline d_i

$$L_{\max} = \max_j (C_j - d_j) \quad \| \max$$

If deadlines are feasible ($L_{\max} < 0$) \rightarrow EDF algo will meet all of them

proof of optimality

Consider order S optimal but not for EDF, then



$$\text{then } \rightarrow C_j = t + p_j, C_k = t + p_j + p_k$$

Consider order S' switching j and k , then

$$C'_k = t + p_k, C'_j = t + p_k + p_j$$

considering $dk < dj$ $C_j = t + p_j \quad C'_j = t + p_k + p_j$

$$C_k = t + p_j + p_k \quad C'_k = t + p_k$$

then: since $d_j > dk$ it's less if subtract

$$C'_j - dj = C_k - dk < C_k - dk \quad \begin{matrix} \text{simply added this} \\ \text{then obviously it's greater} \end{matrix}$$

$$C'_k - dk = t + p_k - dk < t + p_k + p_j - dk = C_k - dk$$

In general $\max_i |C_i - d_i| \leq \max_i |C'_i - d_i| \rightarrow$ switching i and j did not increase $\max_i |C_i - d_i|$
 S' closer to EDF, can be better with different j, k

proof of deadline feasibility

All processes are realizable if $L_{\max} = \max_j (C_j - d_j) < 0$

every deadline is feasible IFF

$$U_i = \frac{1}{d_i} \sum_{k: d_k \leq d_i} P_k \leq 1 \quad \text{with } U = \max_i U_i \quad (\text{load factor})$$

$\frac{1}{d_i}$ · sum of exec
time of processes
before p_i

if $U_i \leq 1 \rightarrow$ all deadline feasible

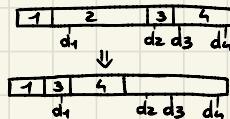
else if $U > 1 \rightarrow$ minimize total of unmet deadlines $\sum U_i$

Moore-Hodgson (MHA) $\rightarrow 1|1|\sum U_i$

tries to minimize the number of tasks that don't meet deadlines $\sum U_i$

- 1) sort tasks j_i based on EDF order
- 2) check if there is task that doesn't meet deadline, if yes continue
- 3) Joc first late task

remove J_B with $P_B = \max_{j \leq i} P_j$
return to 2)



Real time short time scheduling algorithms

release/arrival time

in general $1|r_j|\sum U_i$ is NP-hard

EDF optimal for $1|r_j, p_m t_n|\sum U_i$

$$U_i = \frac{1}{d_i} \sum_{k: d_k \leq d_i} P_k \leq 1 \quad \text{not usable anymore} \Rightarrow U_{[t_1, t_2]} = \frac{1}{t_2 - t_1} \sum_{j: t_1 \leq r_j \leq t_2 \wedge d_j \leq t_2} P_j$$

load $U = \max_i U_i \Rightarrow U = \sup_{t_1 < t_2 < \infty} U_{[t_1, t_2]}$

load $U \leq 1 \Leftrightarrow$ all deadlines feasible

precedence \rightarrow Lawler's algorithm

applicable when precedences in graph $G(V, E)$

optimal for $1|prec|L_{\max}$

- 1) Assume

- G = precedence conditions, $\text{num} = |V|$

- $U(G) = \text{node porzzo}$

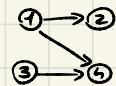
- $T = \sum p_j$ total execution time

- 2) choose task j in $U(G)$ with d_j maximal (or $T-d_j$ minimal)

- task j gets number num

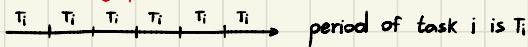
- $\text{num}--$, $T=T-p_j$ remove j 's node from G

- redefine $U(G)$, repeat



$p_1=1 \quad p_2=2 \quad p_3=2 \quad p_4=1$ Proc. times
 $d_1=5 \quad d_2=2 \quad d_3=3 \quad d_4=4$

scheduling periodic tasks



each task n instances with $r_j=jT_i$

instance has relative deadline T_i

EDF can handle n periodic tasks IFF

$$U = \sum_i \frac{p_i}{T_i} \leq 1 \quad (p_i = \text{utilisation of task } i)$$

prove it by showing that $U=U$, EDF optimal for periodic tasks

proof

consider $U = \sup_{t_1 < t_2 < \infty} U[t_1, t_2]$, $U[t_1, t_2] = \frac{1}{t_2 - t_1} \sum_{j: t_1 \leq r_j \leq t_2} p_j$

$$U = \sum_{j=1}^n p_j / T_j \leq 1$$

show that $U \leq U$ for $t_1 < t_2$

it's this
but without LJS, so larger

how many tasks i
can fit in period

$$U[t_1, t_2] (t_2 - t_1) = \sum_{j: t_1 \leq r_j, j \leq t_2} p_j \leq \sum_{j=1}^n \left\lfloor \frac{t_2 - t_1}{T_j} \right\rfloor p_j \leq (t_2 - t_1) U$$

show that $U \geq U$ by replacing $t_1=0$ and $t_2=H$ (H smallest number divisible by T_j)

$$U[0, H] = \sum_{j: r_j \leq H} p_j = \sum_{j=1}^n \frac{H}{T_j} p_j = HU \quad \text{since } U \text{ is the supremum} \rightarrow U \geq U$$

Rate monotonic scheduling

combining periodic real-time tasks with non real-time ones

Handle classical Unix scheduler with added priorities (real time ↑, each one fixed)

organize tasks so that $T_1 \leq T_2 \leq \dots \leq T_n$ more frequent

returns tasks with period T_i priority i (max 1)

RMS respects all deadline certainly if

$$U = \sum_i \frac{p_i}{T_i} \leq n(z^{1/n} - 1) \quad n \rightarrow \infty \quad z^{1/n} - 1 \text{ drops to } \ln z = 0.693$$

RMS optimal in static priority-driven systems (relative deadlines = period)

RMS will meet all deadlines if all periods are multiple of each other and $U \leq 1$